

# DefineAndSolveMLProblem

July 30, 2025

## 1 Lab 8: Define and Solve an ML Problem of Your Choosing

```
[1]: import pandas as pd
import numpy as np
import os
import matplotlib.pyplot as plt
import seaborn as sns
```

In this lab assignment, you will follow the machine learning life cycle and implement a model to solve a machine learning problem of your choosing. You will select a data set and choose a predictive problem that the data set supports. You will then inspect the data with your problem in mind and begin to formulate a project plan. You will then implement the machine learning project plan.

You will complete the following tasks:

1. Build Your DataFrame
2. Define Your ML Problem
3. Perform exploratory data analysis to understand your data.
4. Define Your Project Plan
5. Implement Your Project Plan:
  - Prepare your data for your model.
  - Fit your model to the training data and evaluate your model.
  - Improve your model's performance.

### 1.1 Part 1: Build Your DataFrame

You will have the option to choose one of four data sets that you have worked with in this program:

- The "census" data set that contains Census information from 1994: `censusData.csv`
- Airbnb NYC "listings" data set: `airbnbListingsData.csv`
- World Happiness Report (WHR) data set: `WHR2018Chapter20onlineData.csv`
- Book Review data set: `bookReviewsData.csv`

Note that these are variations of the data sets that you have worked with in this program. For example, some do not include some of the preprocessing necessary for specific models.

**Load a Data Set and Save it as a Pandas DataFrame** The code cell below contains filenames (path + filename) for each of the four data sets available to you.

Task: In the code cell below, use the same method you have been using to load the data using `pd.read_csv()` and save it to DataFrame `df`.

You can load each file as a new DataFrame to inspect the data before choosing your data set.

```
[2]: # File names of the four data sets
adultDataSet_filename = os.path.join(os.getcwd(), "data", "censusData.csv")
airbnbDataSet_filename = os.path.join(os.getcwd(), "data", "airbnbListingsData.
    ↪ csv")
WHRDataSet_filename = os.path.join(os.getcwd(), "data", "WHR2018Chapter2OnlineData.csv")
bookReviewDataSet_filename = os.path.join(os.getcwd(), "data", "bookReviewsData.
    ↪ csv")

df = pd.read_csv(airbnbDataSet_filename, header=0)

df.head()
```

```
[2]:
```

	name \		description \	neighborhood_overview	host_name \		host_location \		host_about	host_response_rate \
0	Skylit Midtown Castle		Beautiful, spacious skylit studio in the heart...	Centrally located in the heart of Manhattan ju...	Jennifer		New York, New York, United States			
1	Whole flr w/private bdrm, bath & kitchen(pls r...		Enjoy 500 s.f. top floor in 1899 brownstone, w...	Just the right mix of urban center and local n...	LisaRoxanne		New York, New York, United States			
2	Spacious Brooklyn Duplex, Patio + Garden		We welcome you to stay in our lovely 2 br dupl...		NaN	Rebecca	Brooklyn, New York, United States			
3	Large Furnished Room Near B'way		Please don't expect the luxury here just a bas...	Theater district, many restaurants around here.	Shunichi		New York, New York, United States			
4	Cozy Clean Guest Room - Family Apt		Our best guests are seeking a safe, clean, spa...	Our neighborhood is full of restaurants and ca...	MaryEllen		New York, New York, United States			

0	A New Yorker since 2000! My passion is creatin...	0.80
1	Laid-back Native New Yorker (formerly bi-coast...	0.09
2	Rebecca is an artist/designer, and Henoch is i...	1.00
3	I used to work for a financial industry but no...	1.00
4	Welcome to family life with my oldest two away...	NaN

	host_acceptance_rate	host_is_superhost	host_listings_count	...	\
0	0.17	True	8.0	...	
1	0.69	True	1.0	...	
2	0.25	True	1.0	...	
3	1.00	True	1.0	...	
4	NaN	True	1.0	...	

	review_scores_communication	review_scores_location	review_scores_value	\
0	4.79	4.86	4.41	
1	4.80	4.71	4.64	
2	5.00	4.50	5.00	
3	4.42	4.87	4.36	
4	4.95	4.94	4.92	

	instant_bookable	calculated_host_listings_count	\
0	False	3	
1	False	1	
2	False	1	
3	False	1	
4	False	1	

	calculated_host_listings_count_entire_homes	\
0	3	
1	1	
2	1	
3	0	
4	0	

	calculated_host_listings_count_private_rooms	\
0	0	
1	0	
2	0	
3	1	
4	1	

	calculated_host_listings_count_shared_rooms	reviews_per_month	\
0	0	0.33	
1	0	4.86	
2	0	0.02	
3	0	3.68	
4	0	0.87	

	n_host_verifications
0	9
1	6
2	3
3	4
4	7

[5 rows x 50 columns]

## 1.2 Part 2: Define Your ML Problem

Next you will formulate your ML Problem. In the markdown cell below, answer the following questions:

1. List the data set you have chosen.
  2. What will you be predicting? What is the label?
  3. Is this a supervised or unsupervised learning problem? Is this a clustering, classification or regression problem? Is it a binary classification or multi-class classification problem?
  4. What are your features? (note: this list may change after you explore your data)
  5. Explain why this is an important problem. In other words, how would a company create value with a model that predicts this label?
1. airbnbDataSet
  2. I will be analyzing which features contribute most to the price level (premium vs budget) of the airbnb. The label will be a "is\_premium" column newly created. I'll be predicting whether an airbnb listing is "premium" or not based on whether its above the threshold (median price).
  3. This is a supervised learning problem. It is a classification problem.
  4. My features are all the other columns except this new "is\_premium" column.
  5. This is an important problem because a "premium" vs "budget" classifier helps hosts and pricing platforms automatically tier listings, recommend competitive pricing strategies, and personalize promotion or search ranking. By understanding which characteristics affect/contribute to premium rates, Airbnb could boost host revenues, improve guest satisfaction with fair pricing, and optimize marketplace liquidity.

## 1.3 Part 3: Understand Your Data

The next step is to perform exploratory data analysis. Inspect and analyze your data set with your machine learning problem in mind. Consider the following as you inspect your data:

1. What data preparation techniques would you like to use? These data preparation techniques may include:
  - addressing missingness, such as replacing missing values with means
  - finding and replacing outliers
  - renaming features and labels
  - finding and replacing outliers

- performing feature engineering techniques such as one-hot encoding on categorical features
  - selecting appropriate features and removing irrelevant features
  - performing specific data cleaning and preprocessing techniques for an NLP problem
  - addressing class imbalance in your data sample to promote fair AI
2. What machine learning model (or models) you would like to use that is suitable for your predictive problem and data?
  - Are there other data preparation techniques that you will need to apply to build a balanced modeling data set for your problem and model? For example, will you need to scale your data?
3. How will you evaluate and improve the model's performance?
  - Are there specific evaluation metrics and methods that are appropriate for your model?

Think of the different techniques you have used to inspect and analyze your data in this course. These include using Pandas to apply data filters, using the Pandas `describe()` method to get insight into key statistics for each column, using the Pandas `dtypes` property to inspect the data type of each column, and using Matplotlib and Seaborn to detect outliers and visualize relationships between features and labels. If you are working on a classification problem, use techniques you have learned to determine if there is class imbalance.

Task: Use the techniques you have learned in this course to inspect and analyze your data. You can import additional packages that you have used in this course that you will need to perform this task.

Note: You can add code cells if needed by going to the Insert menu and clicking on Insert Cell Below in the drop-down menu.

```
[3]: df.shape
```

[3]: (28022, 50)

```
[4]: df['price'].median()
```

```
[4]: 115.0
```

```
[17]: #creating new is_premium column based on the threshold(median)
threshold = df['price'].median()
df['is_premium'] = (df['price'] > threshold).astype(int)
df.head()
```

`Traceback (most recent call last)`

```

~/.local/lib/python3.6/site-packages/pandas/core/indexes/base.py in
↳ get_loc(self, key, method, tolerance)
    2897         try:
-> 2898             return self._engine.get_loc(casted_key)
    2899         except KeyError as err:

```

```

pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_loc()

```

```

pandas/_libs/index.pyx in pandas._libs.index.IndexEngine.get_loc()

```

```

pandas/_libs/hashtable_class_helper.pxi in pandas._libs.hashtable.
↳ PyObjectHashTable.get_item()

```

```

pandas/_libs/hashtable_class_helper.pxi in pandas._libs.hashtable.
↳ PyObjectHashTable.get_item()

```

```

KeyError: 'price'

```

The above exception was the direct cause of the following exception:

```

KeyError                                Traceback (most recent call
↳ last)

```

```

<ipython-input-17-1ebf642a26c6> in <module>()
----> 1 threshold = df['price'].median()
      2 df['is_premium'] = (df['price'] > threshold).astype(int)
      3 df.head()
      4 df["host_response_rate"]

```

```

~/.local/lib/python3.6/site-packages/pandas/core/frame.py in
↳ __getitem__(self, key)
    2904         if self.columns.nlevels > 1:
    2905             return self._getitem_multilevel(key)
-> 2906         indexer = self.columns.get_loc(key)
    2907         if is_integer(indexer):
    2908             indexer = [indexer]

```

```

~/.local/lib/python3.6/site-packages/pandas/core/indexes/base.py in
get_loc(self, key, method, tolerance)
    2898             return self._engine.get_loc(casted_key)
    2899         except KeyError as err:
-> 2900             raise KeyError(key) from err
    2901
    2902         if tolerance is not None:

```

KeyError: 'price'

```

[6]: #checking columns for null values
np.sum(df.isnull(), axis = 0)

```

```

[6]: name                    5
description                  570
neighborhood_overview       9816
host_name                    0
host_location                60
host_about                  10945
host_response_rate          11843
host_acceptance_rate        11113
host_is_superhost            0
host_listings_count          0
host_total_listings_count    0
host_has_profile_pic         0
host_identity_verified       0
neighbourhood_group_cleansed 0
room_type                    0
accommodates                 0
bathrooms                   0
bedrooms                    2918
beds                        1354
amenities                    0
price                        0
minimum_nights               0
maximum_nights               0
minimum_minimum_nights       0
maximum_minimum_nights       0
minimum_maximum_nights       0
maximum_maximum_nights       0
minimum_nights_avg_ntm       0
maximum_nights_avg_ntm       0
has_availability              0
availability_30               0
availability_60               0

```

availability_90	0
availability_365	0
number_of_reviews	0
number_of_reviews_ltm	0
number_of_reviews_l30d	0
review_scores_rating	0
review_scores_cleanliness	0
review_scores_checkin	0
review_scores_communication	0
review_scores_location	0
review_scores_value	0
instant_bookable	0
calculated_host_listings_count	0
calculated_host_listings_count_entire_homes	0
calculated_host_listings_count_private_rooms	0
calculated_host_listings_count_shared_rooms	0
reviews_per_month	0
n_host_verifications	0
is_premium	0
dtype: int64	

```
[7]: df.dtypes
```

name	object
description	object
neighborhood_overview	object
host_name	object
host_location	object
host_about	object
host_response_rate	float64
host_acceptance_rate	float64
host_is_superhost	bool
host_listings_count	float64
host_total_listings_count	float64
host_has_profile_pic	bool
host_identity_verified	bool
neighbourhood_group_cleansed	object
room_type	object
accommodates	int64
bathrooms	float64
bedrooms	float64
beds	float64
amenities	object
price	float64
minimum_nights	int64
maximum_nights	int64
minimum_minimum_nights	float64



maximum_minimum_nights	float64
minimum_maximum_nights	float64
maximum_maximum_nights	float64
minimum_nights_avg_ntm	float64
maximum_nights_avg_ntm	float64
has_availability	bool
availability_30	int64
availability_60	int64
availability_90	int64
availability_365	int64
number_of_reviews	int64
number_of_reviews_ltm	int64
number_of_reviews_l30d	int64
review_scores_rating	float64
review_scores_cleanliness	float64
review_scores_checkin	float64
review_scores_communication	float64
review_scores_location	float64
review_scores_value	float64
instant_bookable	bool
calculated_host_listings_count	int64
calculated_host_listings_count_entire_homes	int64
calculated_host_listings_count_private_rooms	int64
calculated_host_listings_count_shared_rooms	int64
reviews_per_month	float64
n_host_verifications	int64
is_premium	int64
dtype:	object

```
[11]: #dropping columns that aren't relevant/helpful
df.drop(columns = ['price', 'description',
    ↳ 'name', 'neighborhood_overview', 'host_about', 'host_name', 'host_location', 'amenities'],
    ↳ inplace=True)
```

```
[19]: df.drop(columns=['host_response_rate', 'host_acceptance_rate'], inplace=True)
```

```
[18]: df['host_response_rate']
```

```
[18]: 0      0.80
      1      0.09
      2      1.00
      3      1.00
      4      NaN
      ...
      28017    1.00
      28018    0.91
      28019    0.99
```

```
28020    0.90
28021     NaN
Name: host_response_rate, Length: 28022, dtype: float64
```

```
[21]: #replacing null values with the median
df['beds'].fillna(df['beds'].median(), inplace=True)
df['bedrooms'].fillna(df['beds'].median(), inplace=True)
```

```
[22]: np.sum(df.isnull(), axis = 0)
```

```
[22]: host_is_superhost          0
      host_listings_count      0
      host_total_listings_count 0
      host_has_profile_pic      0
      host_identity_verified    0
      neighbourhood_group_cleansed 0
      room_type                 0
      accommodates              0
      bathrooms                 0
      bedrooms                  0
      beds                     0
      minimum_nights            0
      maximum_nights            0
      minimum_minimum_nights    0
      maximum_minimum_nights    0
      minimum_maximum_nights    0
      maximum_maximum_nights    0
      minimum_nights_avg_ntm    0
      maximum_nights_avg_ntm    0
      has_availability           0
      availability_30            0
      availability_60            0
      availability_90            0
      availability_365           0
      number_of_reviews          0
      number_of_reviews_ltm      0
      number_of_reviews_l30d     0
      review_scores_rating       0
      review_scores_cleanliness  0
      review_scores_checkin      0
      review_scores_communication 0
      review_scores_location     0
      review_scores_value        0
      instant_bookable           0
      calculated_host_listings_count 0
      calculated_host_listings_count_entire_homes 0
      calculated_host_listings_count_private_rooms 0
```

```

calculated_host_listings_count_shared_rooms    0
reviews_per_month                             0
n_host_verifications                           0
is_premium                                     0
dtype: int64

```

```

[23]: #one hot encoding columns that have data type of object
to_encode = list(df.select_dtypes(include=['object']).columns)
to_encode
df[to_encode].nunique()

```

```

[23]: neighbourhood_group_cleansed    5
room_type                           4
dtype: int64

```

```

[24]: for colname in to_encode:
        df_encoded = pd.get_dummies(df[colname], prefix=colname + '_')
        df = df.join(df_encoded)

```

```

[25]: df.head()

```

```

[25]:   host_is_superhost  host_listings_count  host_total_listings_count \
0                True                8.0                8.0
1                True                1.0                1.0
2                True                1.0                1.0
3                True                1.0                1.0
4                True                1.0                1.0

   host_has_profile_pic  host_identity_verified  neighbourhood_group_cleansed \
0                True                True                Manhattan
1                True                True                Brooklyn
2                True                True                Brooklyn
3                True                True                Manhattan
4                True                True                Manhattan

   room_type  accommodates  bathrooms  bedrooms  ...  is_premium \
0  Entire home/apt         1         1.0         1.0  ...         1
1  Entire home/apt         3         1.0         1.0  ...         0
2  Entire home/apt         4         1.5         2.0  ...         1
3   Private room          2         1.0         1.0  ...         0
4   Private room          1         1.0         1.0  ...         0

   neighbourhood_group_cleansed__Bronx \
0                0
1                0
2                0
3                0

```

4	0	
---	---	--

  

	neighbourhood_group_cleansed__Brooklyn \
0	0
1	1
2	1
3	0
4	0

  

	neighbourhood_group_cleansed__Manhattan \
0	1
1	0
2	0
3	1
4	1

  

	neighbourhood_group_cleansed__Queens \
0	0
1	0
2	0
3	0
4	0

  

	neighbourhood_group_cleansed__Staten Island	room_type__Entire home/apt \
0	0	1
1	0	1
2	0	1
3	0	0
4	0	0

  

	room_type__Hotel room	room_type__Private room	room_type__Shared room
0	0	0	0
1	0	0	0
2	0	0	0
3	0	1	0
4	0	1	0

[5 rows x 50 columns]

```
[26]: df.drop(columns = to_encode ,axis=1, inplace=True)
```

```
[27]: df.isnull().values.any()
```

```
[27]: False
```

```
[30]: features = list(df.loc[:, df.columns != 'is_premium'])
features
```

```

[30]: ['host_is_superhost',
      'host_listings_count',
      'host_total_listings_count',
      'host_has_profile_pic',
      'host_identity_verified',
      'accommodates',
      'bathrooms',
      'bedrooms',
      'beds',
      'minimum_nights',
      'maximum_nights',
      'minimum_minimum_nights',
      'maximum_minimum_nights',
      'minimum_maximum_nights',
      'maximum_maximum_nights',
      'minimum_nights_avg_ntm',
      'maximum_nights_avg_ntm',
      'has_availability',
      'availability_30',
      'availability_60',
      'availability_90',
      'availability_365',
      'number_of_reviews',
      'number_of_reviews_ltm',
      'number_of_reviews_l30d',
      'review_scores_rating',
      'review_scores_cleanliness',
      'review_scores_checkin',
      'review_scores_communication',
      'review_scores_location',
      'review_scores_value',
      'instant_bookable',
      'calculated_host_listings_count',
      'calculated_host_listings_count_entire_homes',
      'calculated_host_listings_count_private_rooms',
      'calculated_host_listings_count_shared_rooms',
      'reviews_per_month',
      'n_host_verifications',
      'neighbourhood_group_cleansed__Bronx',
      'neighbourhood_group_cleansed__Brooklyn',
      'neighbourhood_group_cleansed__Manhattan',
      'neighbourhood_group_cleansed__Queens',
      'neighbourhood_group_cleansed__Staten Island',
      'room_type__Entire home/apt',
      'room_type__Hotel room',
      'room_type__Private room',
      'room_type__Shared room']

```

## 1.4 Part 4: Define Your Project Plan

Now that you understand your data, in the markdown cell below, define your plan to implement the remaining phases of the machine learning life cycle (data preparation, modeling, evaluation) to solve your ML problem. Answer the following questions:

- Do you have a new feature list? If so, what are the features that you chose to keep and remove after inspecting the data?
- Explain different data preparation techniques that you will use to prepare your data for modeling.
- What is your model (or models)?
- Describe your plan to train your model, analyze its performance and then improve the model. That is, describe your model building, validation and selection plan to produce a model that generalizes well to new data.

Yes there is a new features list in the features variable shown in the cell above. I removed the ones that wouldn't have much significance (including host info), one hot encoded object columns, and for numerical columns I replaced the null values with the median/mean in order to obtain columns without any null values. I will use a DecisionTreeClassifier for this problem. I split the dataset with a 80/20 for the training and test sets. I will also go through different hyperparameters including max\_depth to try to obtain the best/highest accuracy model. Finally, I will review feature importances to see which ones contributed most to the label.

## 1.5 Part 5: Implement Your Project Plan

Task: In the code cell below, import additional packages that you have used in this course that you will need to implement your project plan.

```
[31]: from sklearn.model_selection import train_test_split, GridSearchCV
      from sklearn.tree import DecisionTreeClassifier
      from sklearn.metrics import accuracy_score, roc_auc_score, classification_report
      from sklearn.preprocessing import OneHotEncoder, StandardScaler
      from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
```

Task: Use the rest of this notebook to carry out your project plan.

You will:

1. Prepare your data for your model.
2. Fit your model to the training data and evaluate your model.
3. Improve your model's performance by performing model selection and/or feature selection techniques to find best model for your problem.

Add code cells below and populate the notebook with commentary, code, analyses, results, and figures as you see fit.

```
[34]: y = df['is_premium']
      X = df[features]
```

```
[35]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=123)
```

```
[40]: model = DecisionTreeClassifier(random_state=123)
model.fit(X_train, y_train)
preds = model.predict(X_test)
proba = model.predict_proba(X_test)[: , 1]
print("Accuracy:", accuracy_score(y_test, preds))
print("AUC:",      roc_auc_score(y_test, proba))
```

Accuracy: 0.7637823371989295

AUC: 0.7641754897758394

```
[41]: param_grid = {
    'max_depth': [None, 5, 10, 20],
    'min_samples_leaf': [1, 5, 10],
    'criterion': ['gini', 'entropy']
}
```

```
[42]: grid_dt = GridSearchCV(
    DecisionTreeClassifier(random_state=123),
    param_grid, cv=5, scoring='roc_auc', n_jobs=-1
)
```

```
[43]: grid_dt.fit(X_train, y_train)
best_dt = grid_dt.best_estimator_
print("Best DT params:", grid_dt.best_params_)
print("Best DT CV AUC:", grid_dt.best_score_)
```

Best DT params: {'criterion': 'entropy', 'max\_depth': 10, 'min\_samples\_leaf': 10}

Best DT CV AUC: 0.8873315123579278

```
[45]: importances = pd.Series(best_dt.feature_importances_, index=X_train.columns)
print("\nTop 10 Features:\n", importances.sort_values(ascending=False).head(10))
```

Top 10 Features:

room_type__Entire home/apt	0.436553
neighbourhood_group_cleansed__Manhattan	0.100216
bedrooms	0.051191
accommodates	0.042843
calculated_host_listings_count	0.041407
review_scores_location	0.035859
availability_90	0.019425
availability_60	0.019111
calculated_host_listings_count_private_rooms	0.019012

availability\_365 0.018037  
dtype: float64

```
[46]: feature_imp = best_dt.feature_importances_  
df_features = pd.DataFrame({  
    'feature': X_train.columns,  
    'importance': feature_imp  
})  
df_sorted = df_features.sort_values(by='importance', ascending=False).  
    ↪reset_index(drop=True)  
print("All feature importances:\n", df_sorted)  
print("\nTop 10 features:\n", df_sorted.head(10))
```

All feature importances:

	feature	importance
0	room_type__Entire home/apt	0.436553
1	neighbourhood_group_cleansed__Manhattan	0.100216
2	bedrooms	0.051191
3	accommodates	0.042843
4	calculated_host_listings_count	0.041407
5	review_scores_location	0.035859
6	availability_90	0.019425
7	availability_60	0.019111
8	calculated_host_listings_count_private_rooms	0.019012
9	availability_365	0.018037
10	minimum_nights_avg_ntm	0.017454
11	calculated_host_listings_count_entire_homes	0.016025
12	minimum_minimum_nights	0.014038
13	review_scores_rating	0.013384
14	maximum_minimum_nights	0.012950
15	review_scores_cleanliness	0.012182
16	number_of_reviews	0.012141
17	reviews_per_month	0.011528
18	bathrooms	0.011253
19	neighbourhood_group_cleansed__Brooklyn	0.010705
20	maximum_nights	0.010237
21	n_host_verifications	0.009739
22	review_scores_value	0.008926
23	host_listings_count	0.008383
24	availability_30	0.007475
25	review_scores_checkin	0.005365
26	maximum_maximum_nights	0.005075
27	host_total_listings_count	0.004909
28	review_scores_communication	0.004344
29	beds	0.003495
30	minimum_nights	0.003414
31	number_of_reviews_ltm	0.003083
32	neighbourhood_group_cleansed__Queens	0.002910



33	minimum_maximum_nights	0.002653
34	instant_bookable	0.002087
35	room_type__Hotel room	0.001566
36	maximum_nights_avg_ntm	0.000734
37	number_of_reviews_l30d	0.000292
38	room_type__Private room	0.000000
39	neighbourhood_group_cleansed__Staten Island	0.000000
40	host_is_superhost	0.000000
41	neighbourhood_group_cleansed__Bronx	0.000000
42	calculated_host_listings_count_shared_rooms	0.000000
43	has_availability	0.000000
44	host_identity_verified	0.000000
45	host_has_profile_pic	0.000000
46	room_type__Shared room	0.000000

Top 10 features:

	feature	importance
0	room_type__Entire home/apt	0.436553
1	neighbourhood_group_cleansed__Manhattan	0.100216
2	bedrooms	0.051191
3	accommodates	0.042843
4	calculated_host_listings_count	0.041407
5	review_scores_location	0.035859
6	availability_90	0.019425
7	availability_60	0.019111
8	calculated_host_listings_count_private_rooms	0.019012
9	availability_365	0.018037

From my DecisionTreeClassification modeling, it achieved a value of ROC-AUC of around 0.76. But after tuning my hyperparamters, I found that the best hyperparameters were 'criterion': 'entropy', 'max\_depth': 10, 'min\_samples\_leaf': 10 after using a param\_grid with different hyperparameter values to test out. The AUC value then increased to around 0.887. Then at the end I checked for the features that contributed most (most important) to the label (premium vs budget). The most influential features were listing type (Entire home/apt), location in Manhattan, and number of bedrooms and accommodates.